

# 797 project formulation

*Tobias Kuhlmann, Rui Zhang*

*11/14/2018*

## Data

For our project we simulate univariate data

$$[x_i, y_i] \quad i \in \{1, \dots, n\},$$

where  $y_i \sim iid$  and an unknown and known (different cases) smooth density  $f(x)$ . This may not be just one dataset, but several.

## Models

### Univariate kernel density estimator

We use a univariate kernel density function following Wand and Jones (1995). A density function can be estimated by

$$\hat{f}(x; h) = (nh)^{-1} \sum_{i=1}^n K\{(x - X_i)/h\},$$

where  $K$  is a kernel function satisfying  $\int K(x)dx = 1$  and  $h$  is the bandwidth.

R resources

- Wand and Jones (1995)
- <https://stat.ethz.ch/R-manual/R-devel/library/stats/html/density.html>

### Univariate density estimation with logspline

Let  $B$  be a collection of feasible column vectors following Stone, Hansen, Kooperberg, and Truong (1997). If  $\beta \in B$ , then

$$f(x; \beta) = \exp(\beta_1 B_1(x) + \dots + \beta_J B_J(x) - C(\beta)), L < x < U$$

where

$$C(\beta) = \log\left(\int_L^U \exp(\beta_1 B_1(x) + \dots + \beta_J B_J(x)) dy\right).$$

Then  $f(x; \beta)$  is a positive density function on  $(L, U)$ .

R resources

- <https://www.rdocumentation.org/packages/logspline/versions/2.1.11>
- <https://www.rdocumentation.org/packages/logspline/versions/2.1.11/topics/logspline>
- <https://www.rdocumentation.org/packages/logspline/versions/2.1.11/topics/dlogspline>
- Stone, Hansen, Kooperberg, and Truong (1997)

## Goal

After estimating both models on several sets of simulated data with different sample sizes, our goal is to study and compare the rates of convergence of the MISE as  $n \rightarrow \infty$ .

# Simulation experiment

## Simulation

```
# TODO: Change function to known density to evaluate MISE
# rnorm

# function to estimate
f <- function(x)
  return(exp(3*x)*sin(10*x))

# arbitrary noise level
sigma <- 5
```

## Stats: Univariate Kernel density estimator test

```
y <- rnorm(5000, 0, 1)

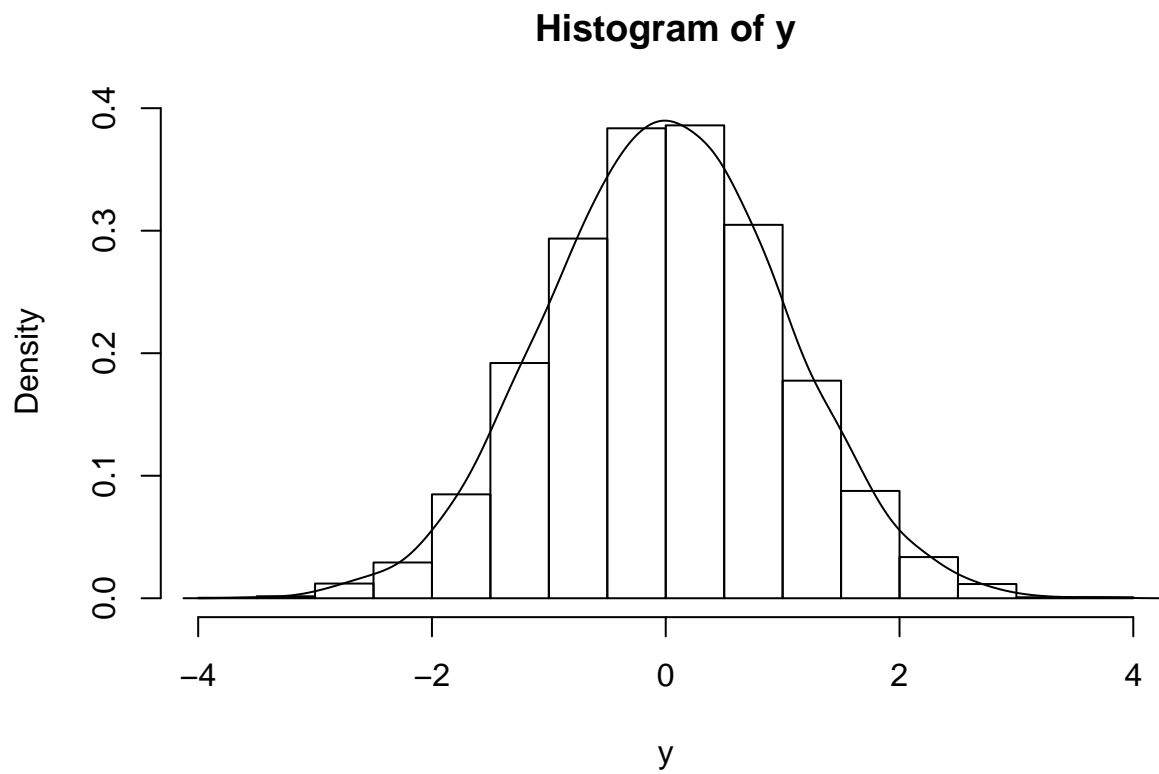
# Univariate kernel density estimator
# use bandwidth estimation as recommended in Venables and Ripley (2002)
fit <- density(y, bw = 'sj')

keep.fits <- fit$y
mise <- mean((fit$y-dnorm(fit$x, 0, 1))^2)
mise

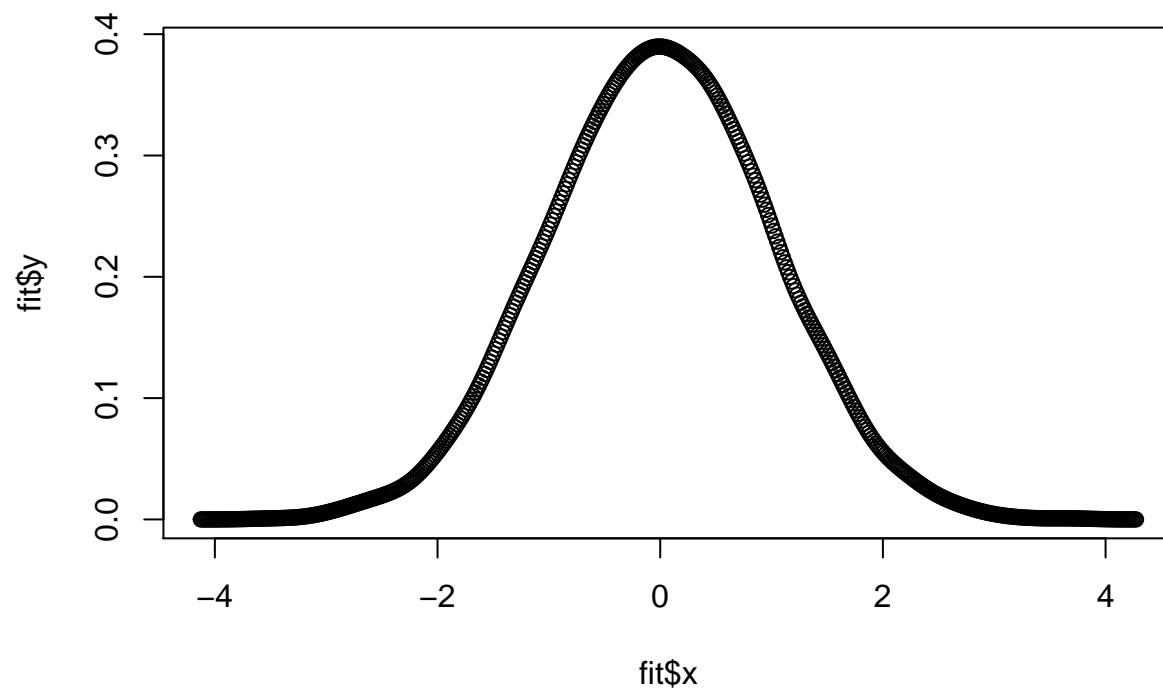
## [1] 1.566223e-05

keep.log.mise <- log10(mise)

# histogram overlay
hist(y, freq = FALSE)
lines(fit)
```



```
# plot fit over x  
plot(x=fit$x, y=fit$y)
```



## KernSmooth (Wand): Univariate Kernel density estimator test

```
# Univariate kernel density estimator following Wand (1995)
#x <- seq(from=-5,to=5,length=1000)
y <- rnorm(1000, 0, 1)

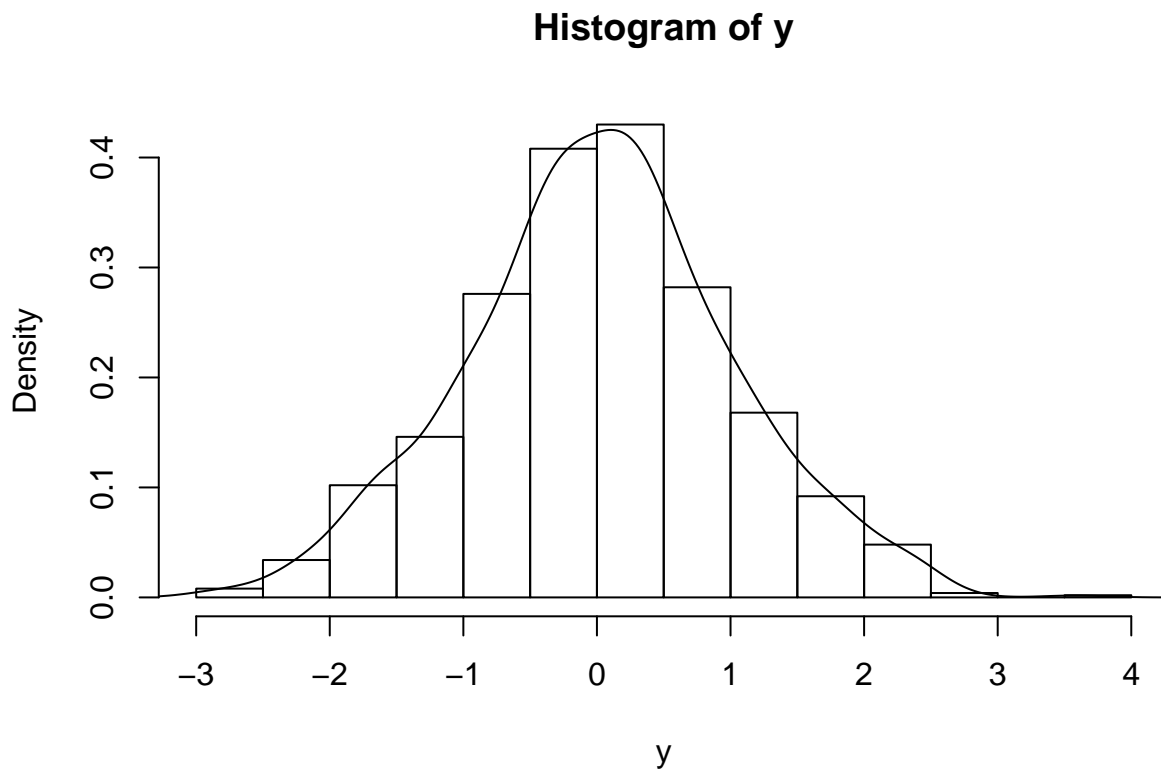
# select optimal bandwidth
h <- dpik(y)
# kde following Wand (1995)
fit <- bkde(x=y, bandwidth=h)

keep.fits <- fit$y
mise <- mean((fit$y-dnorm(fit$x, 0, 1))^2)
mise
```

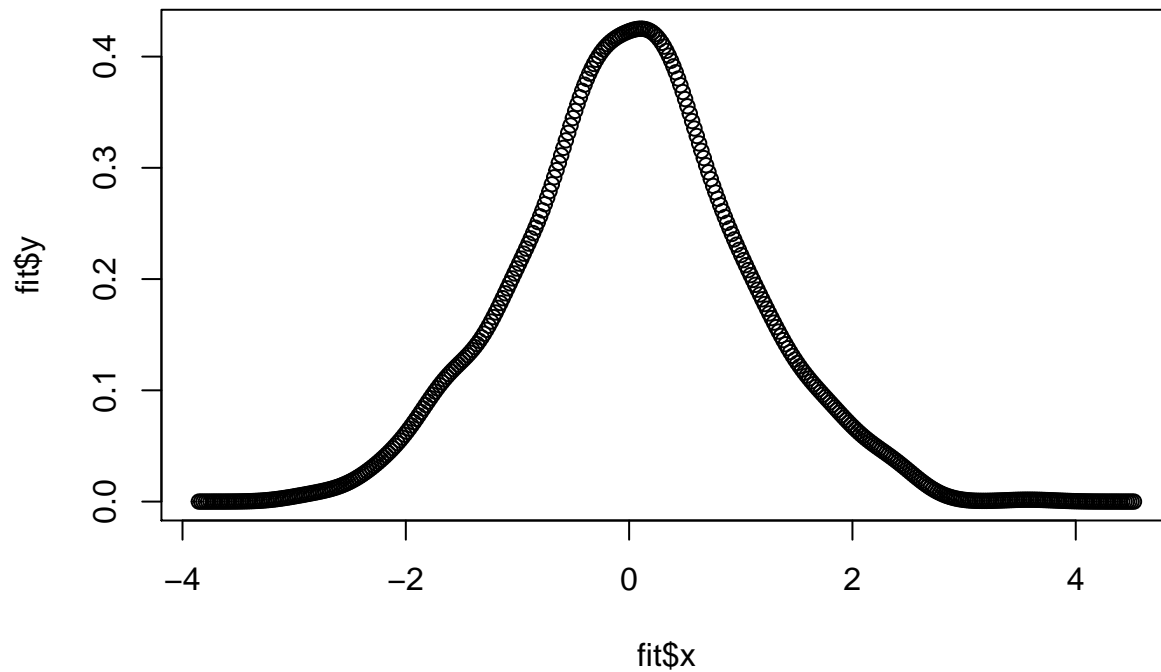
```
## [1] 0.0001844793
```

```
keep.log.mise <- log10(mise)
```

```
# histogram overlay
hist(y, freq = FALSE)
lines(fit)
```



```
# plot fit over x
plot(x=fit$x, y=fit$y)
```



### Logsplines density estimator test

```

y <- rnorm(5000, 0, 1)

# logspline density estimator
fit <- logspline(y)
# summary(fit)
# density object
x = seq(from=-5, to=5, length.out=401)
dens <- dlogspline(q=x, fit=fit)
#summary(dens)

# MISE: mean(dlogspline quantiles - true_quantiles)^2
keep.fits <- dens
mise <- mean((dens-dnorm(x, 0, 1))^2)
mise

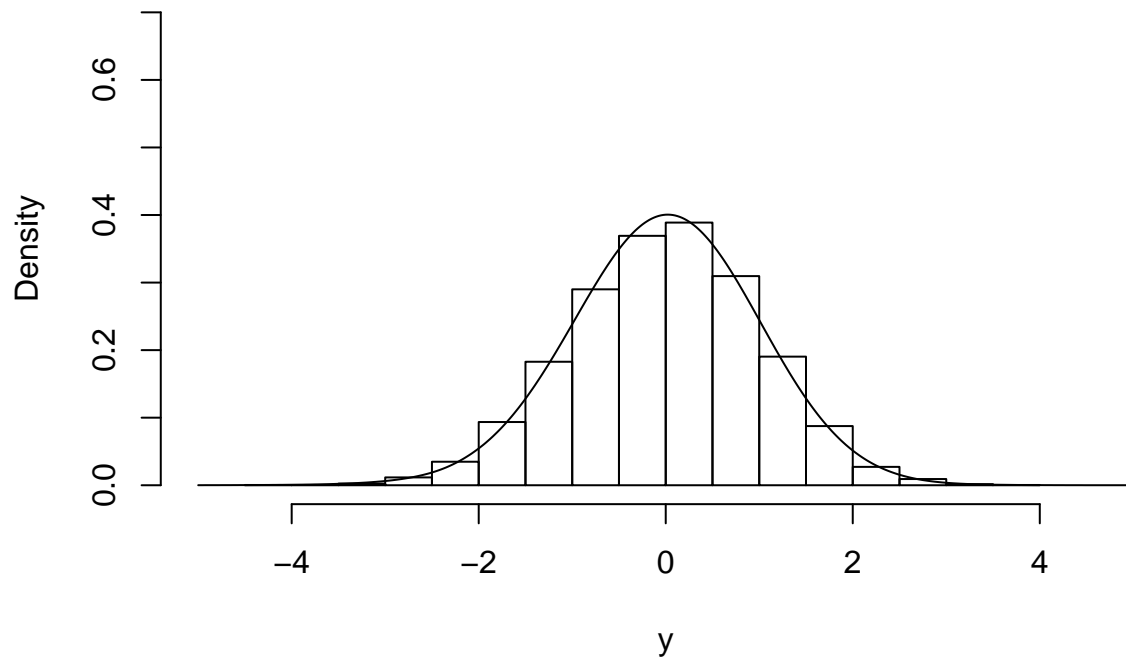
## [1] 3.500502e-06

keep.log.mise <- log10(mise)

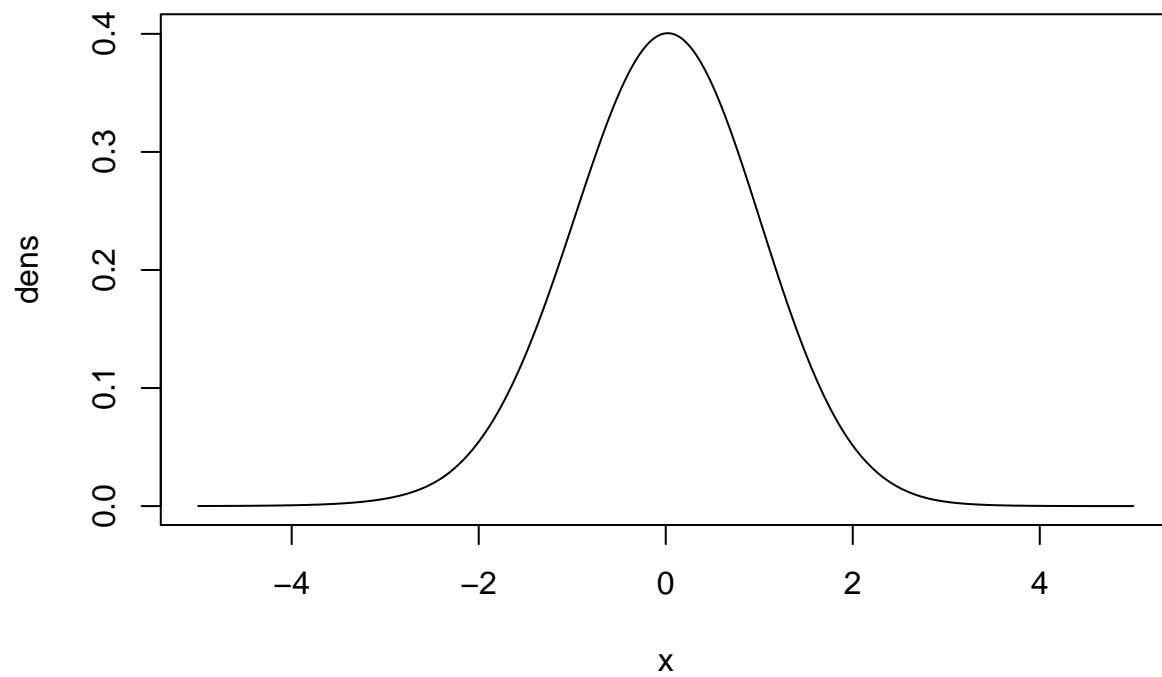
# histogram overlay
hist(y, freq = FALSE, xlim=c(-5,5), ylim=c(0,0.7))
# plot density of logsplinefit
#plot(fit, n = 101, what = "d")
# density overlay
lines(x, dens, type = "l")

```

## Histogram of y



```
# plot dlogspline fit over x  
plot(x, dens, type = "l")
```



Monte Carlo experiment

Normal distribution

```

reps.per.n <- 10
log10.ns <- seq(from=2,to=5,length=20) # equally space n's on log scale
ns <- round(10^log10.ns)
log10.ns <- log10(ns)

# make storage for what we want to keep
keep.kernel.wand <- data.frame(log.n = rep(log10.ns,each=reps.per.n),log.mise=NA)
keep.kernel <- data.frame(log.n = rep(log10.ns,each=reps.per.n),log.mise=NA)
keep.logspline <- data.frame(log.n = rep(log10.ns,each=reps.per.n),log.mise=NA)

# let's keep the fits too (it's always useful to look at estimates)
keep.kernel.wand.fits <- matrix(NA,length(keep.kernel$log.n),401)
keep.kernel.fits <- matrix(NA,length(keep.kernel$log.n),401)
keep.logspline.fits <- matrix(NA,length(keep.logspline$log.n),401)

counter <- 1
for (n.i in ns)
{
  for (mc.i in 1:reps.per.n)
  {
    # generate data
    # TODO1: add some noise to distributions?
    y <- rnorm(n.i, 0, 1)

    # Univariate kernel density estimator from KernSmooth package (Wand (1995))
    # select optimal bandwidth
    h <- dpik(y)
    # kde following Wand (1995)
    fit <- bkde(x=y, bandwidth=h, gridsize = 401L)
    # keep fits
    keep.kernel.wand.fits[counter,] <- fit$y
    # calc mse
    mise <- mean((fit$y-dnorm(fit$x, 0, 1))^2)
    # log mise
    keep.kernel.wand$log.mise[counter] <- log10(mise)

    # store kde x values for log spline evaluation
    x = fit$x

    # Univariate kernel density estimator from stats package
    # use smoothing parameter (standard deviation of kernel) 'sj', as recommended in Venables and Ripley
    fit <- density(y, bw = 'sj', n=401)
    # keep fits
    keep.kernel.fits[counter,] <- fit$y
    # calc mse
    mise <- mean((fit$y-dnorm(fit$x, 0, 1))^2)
    # log mise
    keep.kernel$log.mise[counter] <- log10(mise)

    # Logspline density estimator

```

```

fit <- logspline(y)
# density values on 401 equally spaced points
dens <- dlogspline(q=x, fit=fit)
# keep fits
keep.logspline.fits[counter,] <- dens
# calc mise
mise <- mean((dens-dnorm(x, 0, 1))^2)
# log mise
keep.logspline$log.mise[counter] <- log10(mise)

# runtime
counter <- counter + 1
if (counter %% 20 == 0){
  print(paste0("Iteration: ", counter))
}
}

```

```

## [1] "Iteration: 20"
## [1] "Iteration: 40"
## [1] "Iteration: 60"
## [1] "Iteration: 80"
## [1] "Iteration: 100"
## [1] "Iteration: 120"
## [1] "Iteration: 140"
## [1] "Iteration: 160"
## [1] "Iteration: 180"
## [1] "Iteration: 200"

```

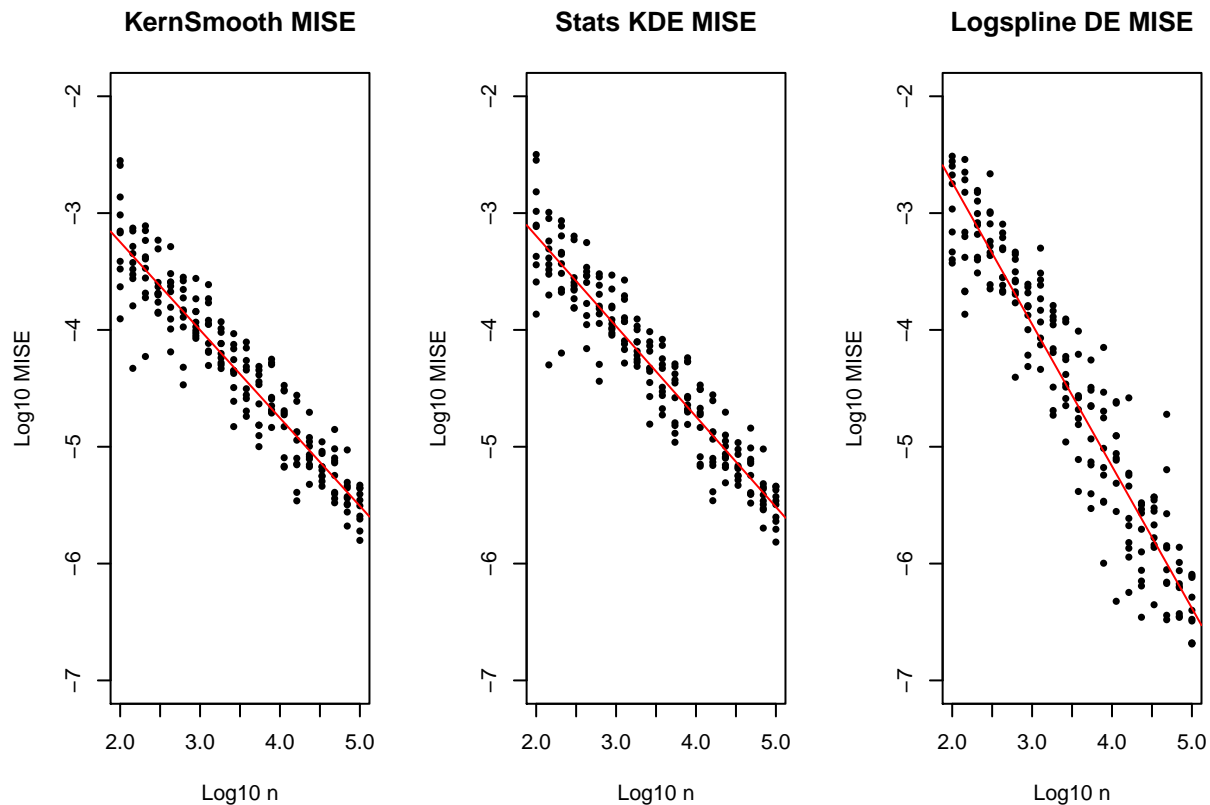
## Visualization

```

# plot results
par(mfrow=c(1,3))
# kde Wand (1995) plot
plot(keep.kernel.wand$log.n,keep.kernel.wand$log.mise,xlab="Log10 n",ylab="Log10 MISE",type="n", ylim=c(
points(keep.kernel.wand$log.n,keep.kernel.wand$log.mise,pch=16,cex=.75)
lmfit_kernel.wand <- lm(log.mise~log.n,data=keep.kernel.wand)
abline(lmfit_kernel.wand,col="red")
title("KernSmooth MISE")
# kde plot
plot(keep.kernel$log.n,keep.kernel$log.mise,xlab="Log10 n",ylab="Log10 MISE",type="n", ylim=c(-7,-2))
points(keep.kernel$log.n,keep.kernel$log.mise,pch=16,cex=.75)
lmfit_kernel <- lm(log.mise~log.n,data=keep.kernel)
abline(lmfit_kernel,col="red")
title("Stats KDE MISE")
# logspline plot
plot(keep.logspline$log.n,keep.logspline$log.mise,xlab="Log10 n",ylab="Log10 MISE",type="n", ylim=c(-7,
points(keep.logspline$log.n,keep.logspline$log.mise,pch=16,cex=.75)
lmfit_logspline <- lm(log.mise~log.n,data=keep.logspline)
abline(lmfit_logspline,col="red")
title("Logspline DE MISE")

```





```
# look at kernel coefficient of interest
print(summary(lmfit_kernel.wand)$coefficients)
```

```
##              Estimate Std. Error  t value    Pr(>|t|)
## (Intercept) -1.7413043 0.06903777 -25.22249 9.477074e-64
## log.n       -0.7532185 0.01908953 -39.45715 9.299547e-96
```

```
print(confint(lmfit_kernel.wand))
```

```
##              2.5 %    97.5 %
## (Intercept) -1.8774480 -1.6051606
## log.n       -0.7908634 -0.7155737
```

```
# look at kernel coefficient of interest
print(summary(lmfit_kernel)$coefficients)
```

```
##              Estimate Std. Error  t value    Pr(>|t|)
## (Intercept) -1.6488327 0.06968400 -23.66157 1.282456e-59
## log.n       -0.7732373 0.01926822 -40.13018 4.730646e-97
```

```
print(confint(lmfit_kernel))
```

```
##              2.5 %    97.5 %
## (Intercept) -1.7862507 -1.511415
## log.n       -0.8112345 -0.735240
```

```
# look at coefficient of interest
print(summary(lmfit_logspline)$coefficients)
```

```
##              Estimate Std. Error  t value    Pr(>|t|)
## (Intercept) -0.3063016 0.11115880 -2.755532 6.406512e-03
## log.n       -1.2151070 0.03073636 -39.533214 6.628501e-96
```

```

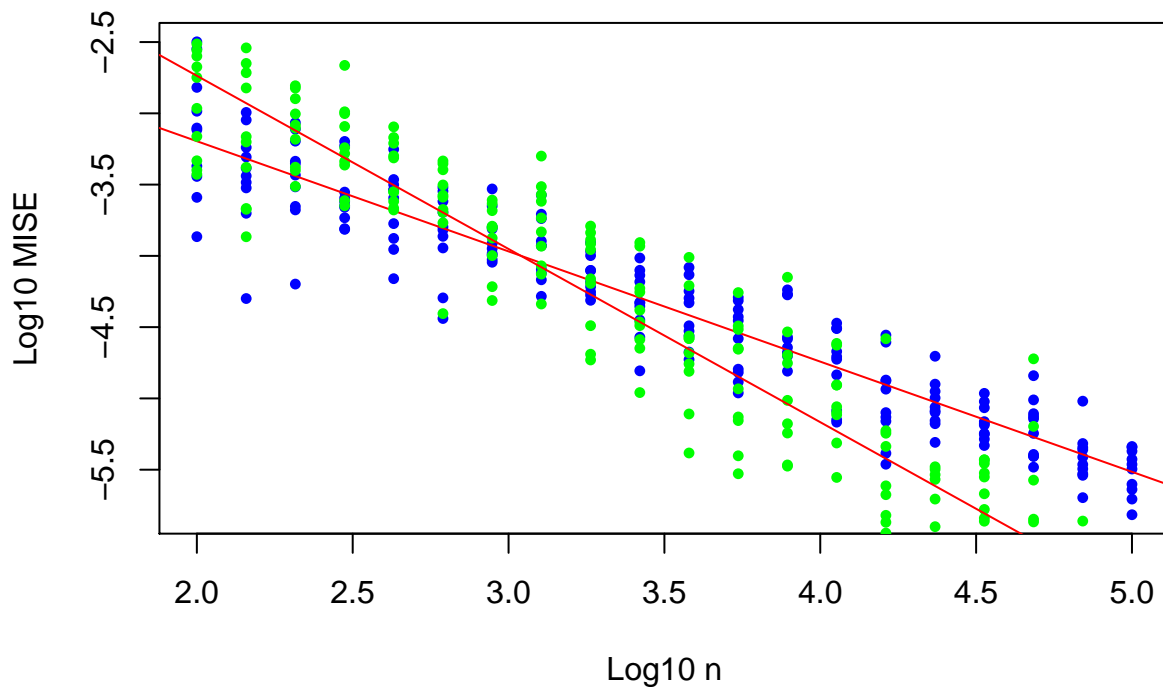
print(confint(lmfit_logspline))

##                2.5 %      97.5 %
## (Intercept) -0.5255087 -0.0870945
## log.n       -1.2757196 -1.1544944

# plot results
par(mfrow=c(1,1))
# kernel
plot(keep.kernel$log.n,keep.kernel$log.mise,xlab="Log10 n",ylab="Log10 MISE",type="n")
points(keep.kernel$log.n,keep.kernel$log.mise,pch=16,cex=.75, col='blue')
lmfit <- lm(log.mise~log.n,data=keep.kernel)
abline(lmfit,col="red")
# logspline
points(keep.logspline$log.n,keep.logspline$log.mise,pch=16,cex=.75, col='green')
lmfit <- lm(log.mise~log.n,data=keep.logspline)
abline(lmfit,col="red")
title("Density estimation asymptotic MISE")

```

## Density estimation asymptotic MISE



```

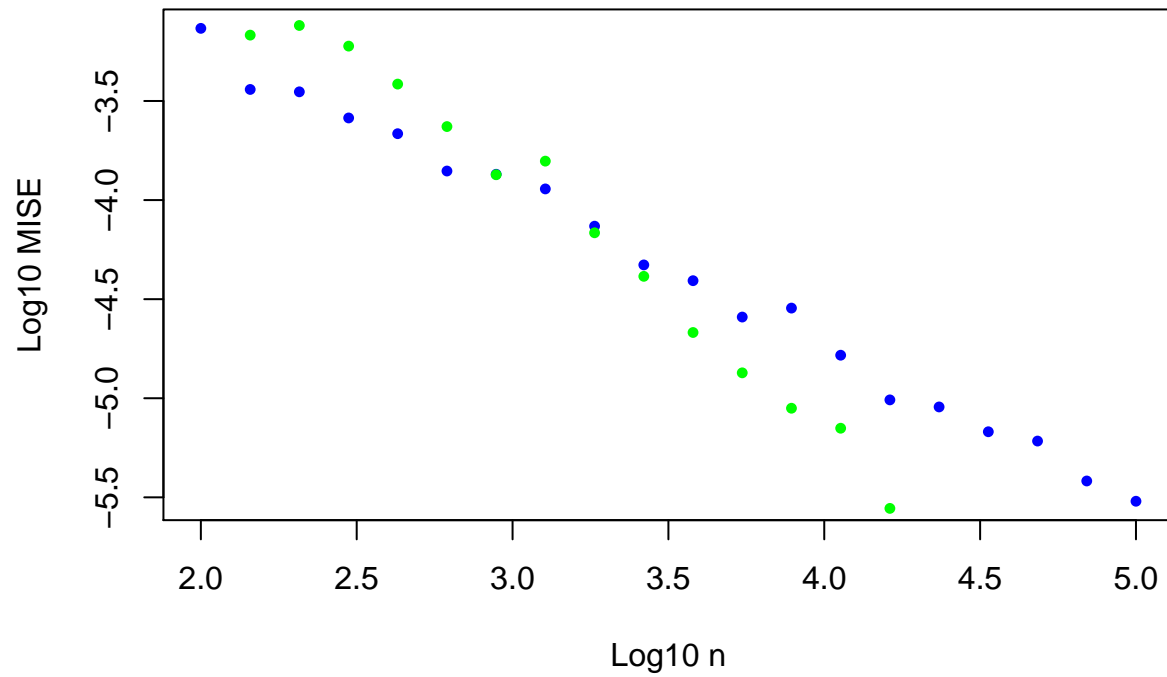
# legend

# mean plots
kernel_mean_mise = aggregate(keep.kernel, by=list(keep.kernel$log.n), mean)
logspline_mean_mise = aggregate(keep.logspline, by=list(keep.logspline$log.n), mean)

plot(kernel_mean_mise$log.n,kernel_mean_mise$log.mise,xlab="Log10 n",ylab="Log10 MISE",type="n")
points(kernel_mean_mise$log.n,kernel_mean_mise$log.mise,pch=16,cex=.75, col='blue')
points(logspline_mean_mise$log.n,logspline_mean_mise$log.mise,pch=16,cex=.75, col='green')
title("Density estimation asymptotic MISE")

```

## Density estimation asymptotic MISE



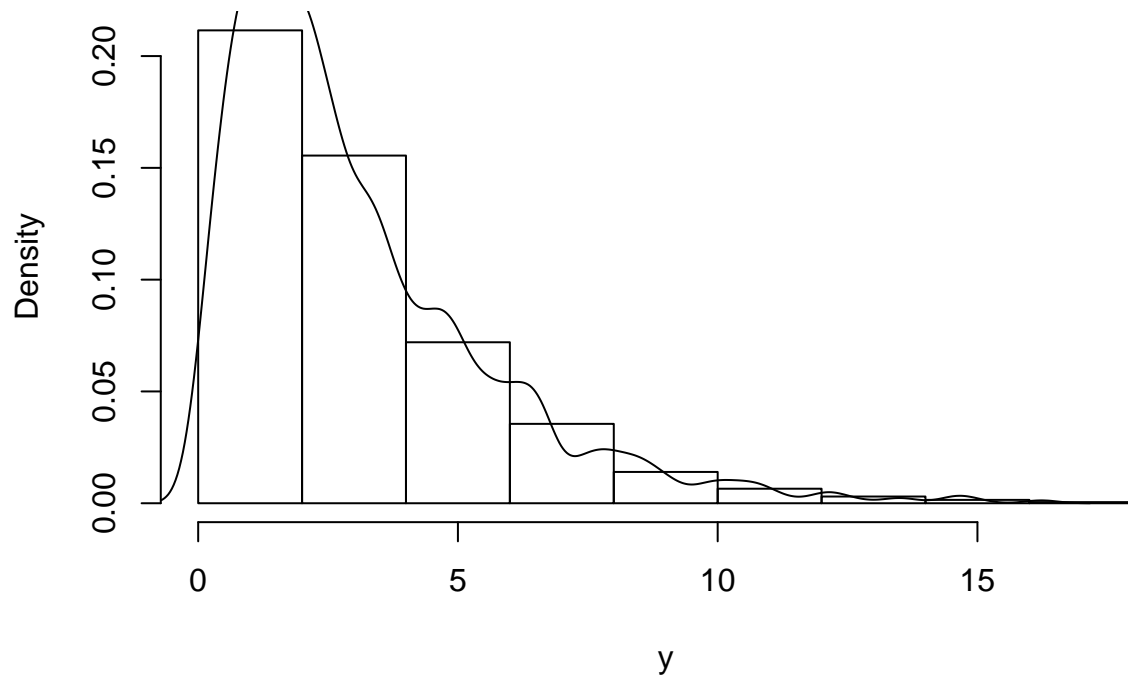
## Chi squared distribution

```
y <- rchisq(seq(from=0,to=25,length=1000), df=3)

# Univariate kernel density estimator
# use bandwidth estimation as recommended in Venables and Ripley (2002)
fit <- density(y, bw = 'sj')

# histogram overlay
hist(y, freq = FALSE)
lines(fit)
```

## Histogram of y

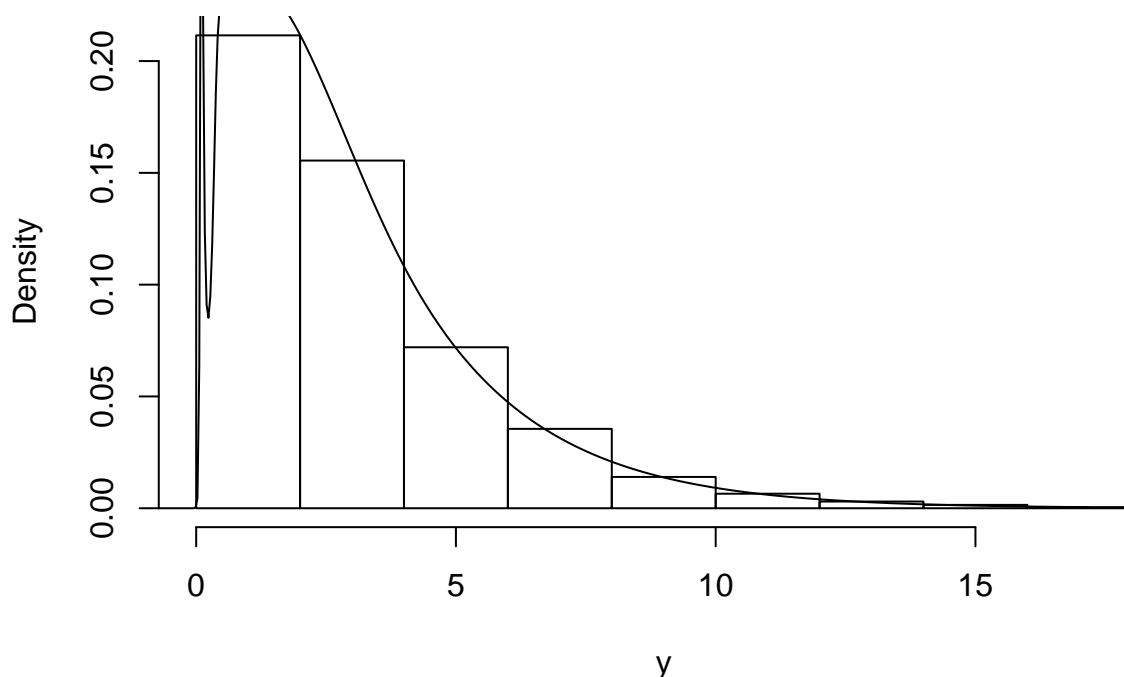


```
x <- fit$x

# logspline density estimator
fit <- logspline(y)
# summary(fit)
# density object
dens <- dlogspline(q=x, fit=fit)
#summary(dens)

# histogram overlay
hist(y, freq = FALSE)
# plot density of logsplinefit
#plot(fit, n = 101, what = "d")
# density overlay
lines(x, dens, type = "l")
```

## Histogram of y



```

reps.per.n <- 10
log10.ns <- seq(from=2,to=5,length=20) # equally space n's on log scale
ns <- round(10^log10.ns)
log10.ns <- log10(ns)

# make storage for what we want to keep
keep.kernel.wand <- data.frame(log.n = rep(log10.ns,each=reps.per.n),log.mise=NA)
keep.kernel <- data.frame(log.n = rep(log10.ns,each=reps.per.n),log.mise=NA)
keep.logspline <- data.frame(log.n = rep(log10.ns,each=reps.per.n),log.mise=NA)

# let's keep the fits too (it's always useful to look at estimates)
keep.kernel.wand.fits <- matrix(NA,length(keep.kernel$log.n),401)
keep.kernel.fits <- matrix(NA,length(keep.kernel$log.n),401)
keep.logspline.fits <- matrix(NA,length(keep.logspline$log.n),401)

counter <- 1
for (n.i in ns)
{
  for (mc.i in 1:reps.per.n)
  {
    # generate data
    # TODO1: add some noise to distributions!
    y <- rchisq(n.i, df=3)

    # Univariate kernel density estimator from KernSmooth package (Wand (1995))
    # select optimal bandwidth
    h <- dpik(y)
    # kde following Wand (1995)
    fit <- bkde(x=y, bandwidth=h, kernel='normal', gridsize = 401L)
  }
  counter <- counter + 1
}

```

```

# keep fits
keep.kernel.wand.fits[counter,] <- fit$y
# calc mse
mise <- mean((fit$y-dnorm(fit$x, 0, 1))^2)
# log mise
keep.kernel.wand$log.mise[counter] <- log10(mise)

# store kde x values for log spline evaluation
x = fit$x

# Univariate kernel density estimator from stats package
# use smoothing parameter (standard deviation of kernel) 'sj', as recommended in Venables and Ripley
fit <- density(y, bw = 'sj', n=401)
# keep fits
keep.kernel.fits[counter,] <- fit$y
# calc mse
mise <- mean((fit$y-dnorm(fit$x, 0, 1))^2)
# log mise
keep.kernel$log.mise[counter] <- log10(mise)

# Logspline density estimator
fit <- logspline(y)
# density values on 401 equally spaced points
dens <- dlogspline(q=x, fit=fit)
# keep fits
keep.logspline.fits[counter,] <- dens
# calc mise
mise <- mean((dens-dchisq(x, df=3))^2)
# log mise
keep.logspline$log.mise[counter] <- log10(mise)

# runtime
counter <- counter + 1
if (counter %% 20 == 0){
  print(paste0("Iteration: ", counter))
}
}
}

```

```
## [1] "Iteration: 20"
```

```
## [1] "Iteration: 40"
```

```
## Warning in logspline(y): Not all models could be fitted
```

```
## [1] "Iteration: 60"
```

```
## [1] "Iteration: 80"
```

```
## [1] "Iteration: 100"
```

```
## [1] "Iteration: 120"
```

```
## [1] "Iteration: 140"
```

```
## [1] "Iteration: 160"
```

```
## [1] "Iteration: 180"
```

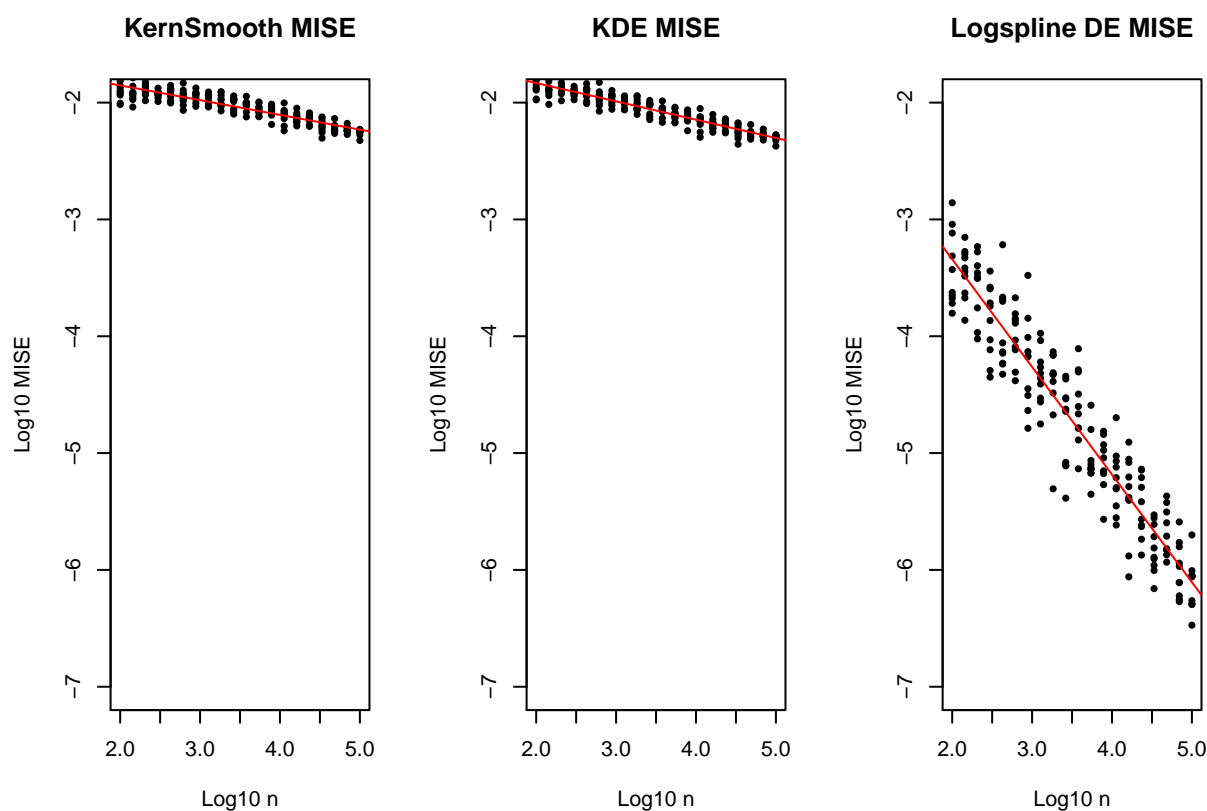
```
## [1] "Iteration: 200"
```

## Visualization

```
# plot results
# plot results
par(mfrow=c(1,3))
# kde Wand (1995) plot
plot(keep.kernel.wand$log.n,keep.kernel.wand$log.mise,xlab="Log10 n",ylab="Log10 MISE",type="n", ylim=c(-7,-2))
points(keep.kernel.wand$log.n,keep.kernel.wand$log.mise,pch=16,cex=.75)
lmfit_kernel.wand <- lm(log.mise~log.n,data=keep.kernel.wand)
abline(lmfit_kernel.wand,col="red")
title("KernSmooth MISE")

# stat kde
plot(keep.kernel$log.n,keep.kernel$log.mise,xlab="Log10 n",ylab="Log10 MISE",type="n", ylim=c(-7,-2))
points(keep.kernel$log.n,keep.kernel$log.mise,pch=16,cex=.75)
lmfit_kernel <- lm(log.mise~log.n,data=keep.kernel)
abline(lmfit_kernel,col="red")
title("KDE MISE")

# logspline de
plot(keep.logspline$log.n,keep.logspline$log.mise,xlab="Log10 n",ylab="Log10 MISE",type="n", ylim=c(-7,-2))
points(keep.logspline$log.n,keep.logspline$log.mise,pch=16,cex=.75)
lmfit_logspline <- lm(log.mise~log.n,data=keep.logspline)
abline(lmfit_logspline,col="red")
title("Logspline DE MISE")
```



```
# look at kernel coefficient of interest
print(summary(lmfit_kernel)$coefficients)
```

```
##               Estimate Std. Error   t value    Pr(>|t|)
## (Intercept) -1.518808 0.014600015 -104.02786 8.946046e-175
## log.n       -0.156863 0.004037029  -38.85604 1.375058e-94

print(confint(lmfit_kernel))

##               2.5 %    97.5 %
## (Intercept) -1.5475998 -1.4900168
## log.n       -0.1648241 -0.1489019

# look at coefficient of interest
print(summary(lmfit_logspline)$coefficients)

##               Estimate Std. Error   t value    Pr(>|t|)
## (Intercept) -1.4994096 0.07974696  -18.80209 6.365343e-46
## log.n       -0.9211373 0.02205072  -41.77358 3.859145e-100

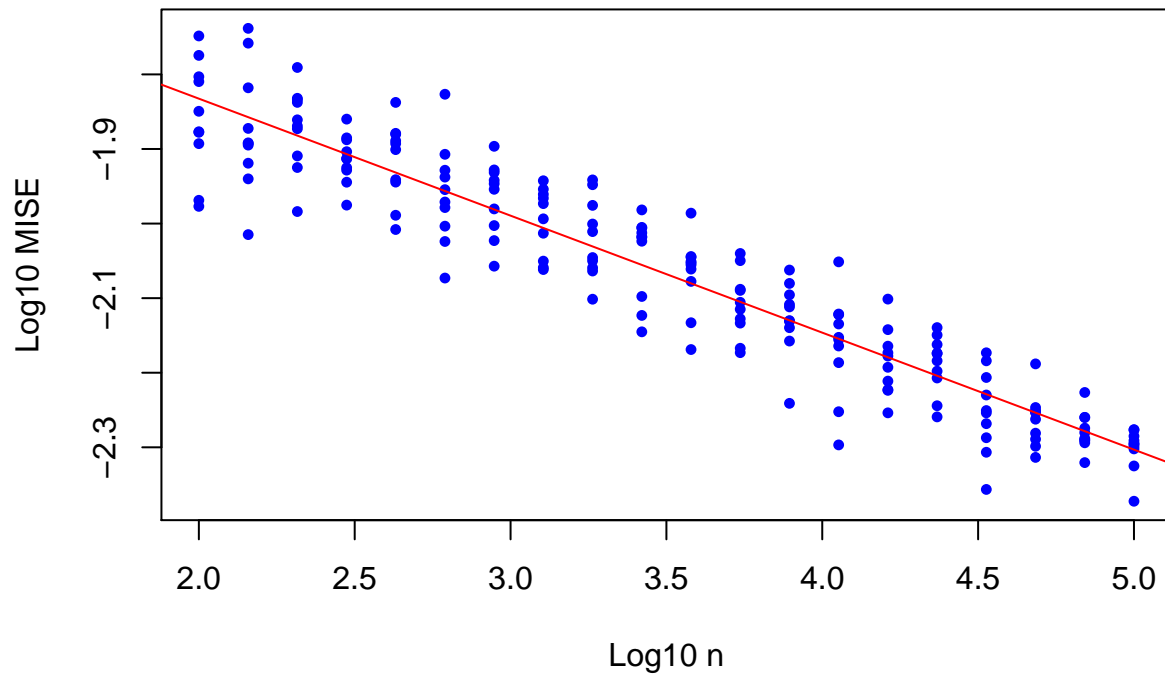
print(confint(lmfit_logspline))

##               2.5 %    97.5 %
## (Intercept) -1.6566720 -1.3421472
## log.n       -0.9646217 -0.8776529

# plot results
par(mfrow=c(1,1)) # Assumes 20 different sample sizes
# kernel
plot(keep.kernel$log.n,keep.kernel$log.mise,xlab="Log10 n",ylab="Log10 MISE",type="n")
points(keep.kernel$log.n,keep.kernel$log.mise,pch=16,cex=.75, col='blue')
lmfit <- lm(log.mise~log.n,data=keep.kernel)
abline(lmfit,col="red")
# logspline
points(keep.logspline$log.n,keep.logspline$log.mise,pch=16,cex=.75, col='green')
lmfit <- lm(log.mise~log.n,data=keep.logspline)
abline(lmfit,col="red")
title("Density estimation asymptotic MISE")
```



## Density estimation asymptotic MISE

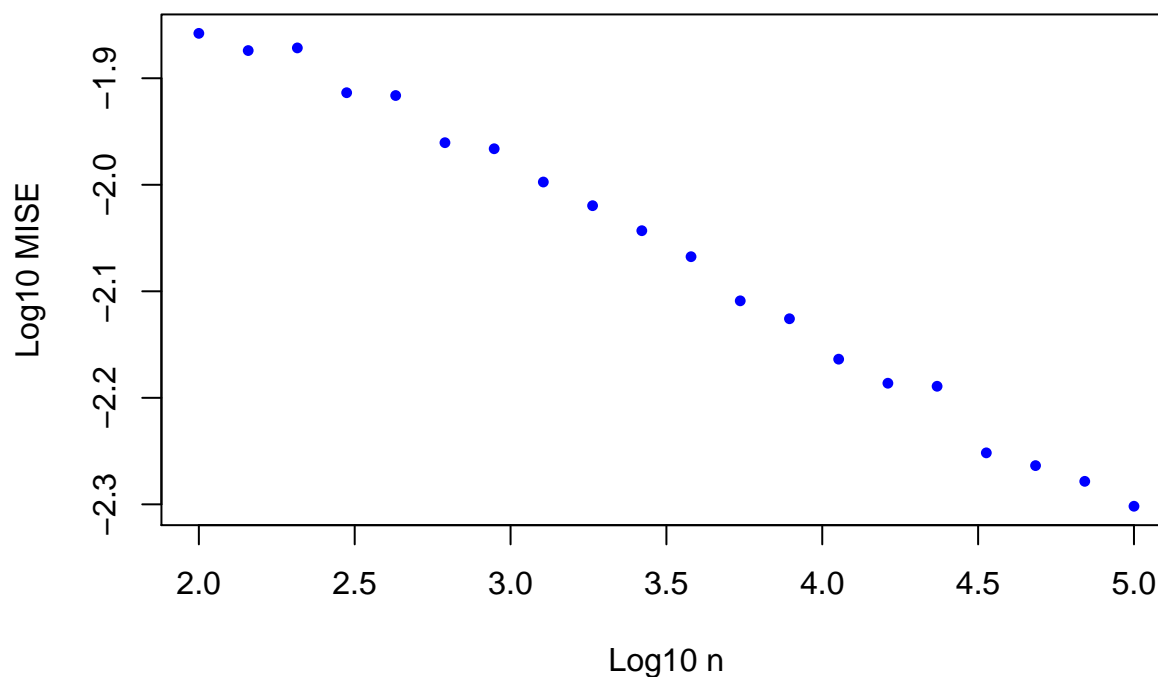


```
# legend

# mean plots
kernel_mean_mise = aggregate(keep.kernel, by=list(keep.kernel$log.n), mean)
logspline_mean_mise = aggregate(keep.logspline, by=list(keep.logspline$log.n), mean)

plot(kernel_mean_mise$log.n, kernel_mean_mise$log.mise, xlab="Log10 n", ylab="Log10 MISE", type="n")
points(kernel_mean_mise$log.n, kernel_mean_mise$log.mise, pch=16, cex=.75, col='blue')
points(logspline_mean_mise$log.n, logspline_mean_mise$log.mise, pch=16, cex=.75, col='green')
title("Density estimation asymptotic MISE")
```

## Density estimation asymptotic MISE



### Questions

- Why are convergence rates different (normal vs. chi2)? Maybe depend on Kernel?

### Conclusions

- Is logspline log mise asymptotic behaviour linear?
- Logspline log mise goes faster to zero than kernel density estimation!

```
#par(mfrow=c(4,5),mar=c(1,1,1,1)) # Assumes 20 different sample sizes
#for (log.n.i in unique(keep$log.n))
#{
#   junk <- keep.fits[keep$log.n==log.n.i,]
#   plot(fit$x,f(fit$x),type="n",ylim=range(junk),main=paste("n=",10^log.n.i),
#   xlab="n",ylab="n",axes=F)
#   for (i in 1:dim(junk)[1])
#     lines(fit$x,junk[i,])
#   lines(fit$x,f(fit$x),col="red")
#}
#}
```